



Scalable Deep Learning Using MXNet

Alex Smola

AWS Machine Learning

Why yet another deep networks tool?

Caffe
Torch
Theano
Tensorflow
CNTK
Keras
Paddle
Chainer
SINGA
DL4J

image credit - Banksy/wikipedia

Why yet another deep networks tool?

- **Frugality & resource efficiency**


Engineered for cheap GPUs with smaller memory, slow networks

- **Speed**

- Linear scaling with `#machines` and `#GPUs`
- High efficiency on single machine, too (C++ backend)

- **Simplicity**

Mix declarative and imperative code



single implementation of
backend system and
common operators

performance guarantee
regardless which frontend
language is used

frontend

backend

Imperative Programs



```
import numpy as np
a = np.ones(10)
b = np.ones(10) * 2
c = b * a
print c
d = c + 1
```

Easy to tweak
with python
codes


Pro

- Straightforward and flexible.
- Take advantage of language native features (loop, condition, debugger)

Con

- Hard to optimize

Declarative Programs



```
A = Variable('A')
B = Variable('B')
C = B * A
D = C + 1
f = compile(D)
d = f(A=np.ones(10),
      B=np.ones(10)*2)
```

Pro

- More chances for optimization
- Cross different languages


Con

- Less flexible

C can share memory with D,
because C is deleted later

Imperative vs. Declarative for Deep Learning

Computational Graph
of the Deep Architecture




Needs heavy optimization,
fits **declarative** programs

Updates and Interactions
with the graph

- Iteration loops
- Parameter update
$$w \leftarrow w - \eta \partial_w f(w)$$
- Beam search
- Feature extraction ...

Needs mutation and more
language native features, good for
imperative programs

LeNet in R (using MXNet)

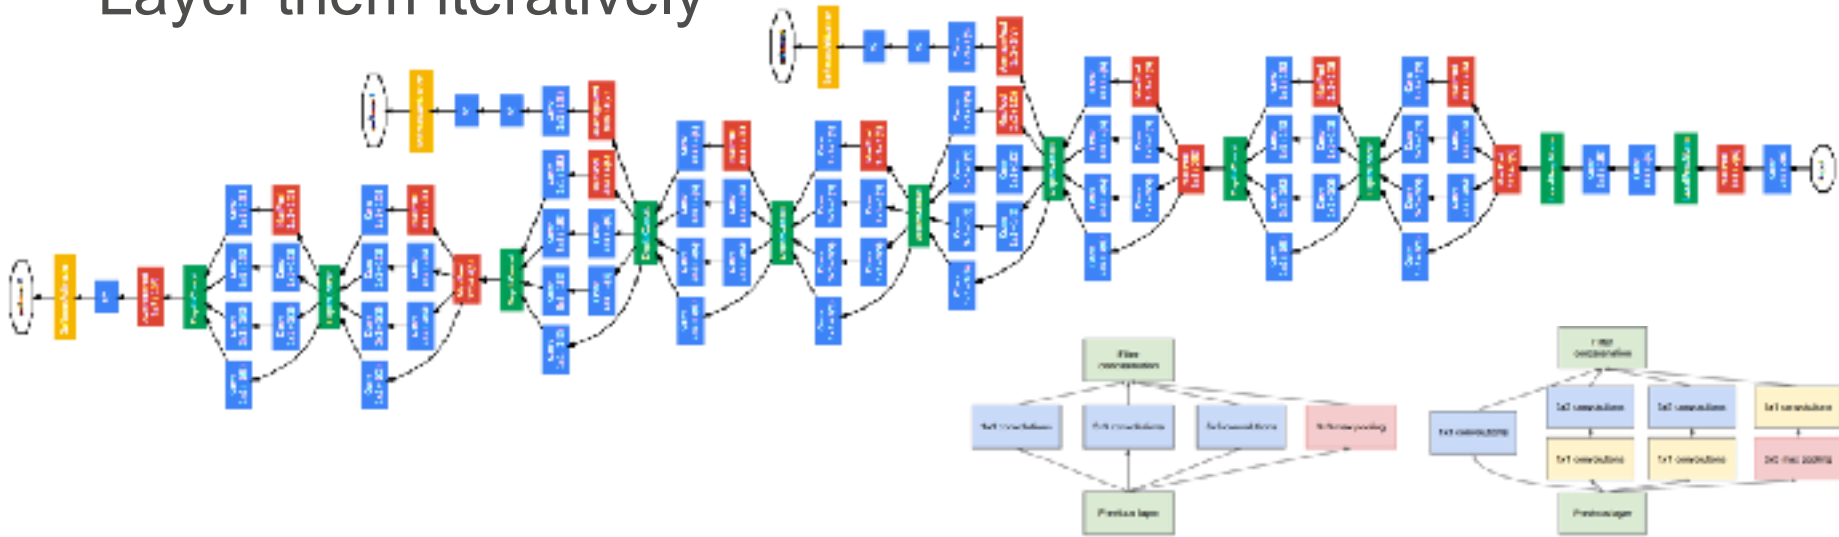


LeNet ~1995

```
get_symbol <- function(num_classes = 1000) {  
  data <- mx.symbol.Variable('data')  
  conv1 <- mx.symbol.Convolution(data = data, kernel = c(5, 5), num_filter = 20)  
  tanh1 <- mx.symbol.Activation(data = conv1, act_type = "tanh")  
  pool1 <- mx.symbol.Pooling(data = tanh1, pool_type = "max", kernel = c(2, 2), stride = c(2, 2))  
  conv2 <- mx.symbol.Convolution(data = pool1, kernel = c(5, 5), num_filter = 50)  
  tanh2 <- mx.symbol.Activation(data = conv2, act_type = "tanh")  
  pool2 <- mx.symbol.Pooling(data = tanh2, pool_type = "max", kernel = c(2, 2), stride = c(2, 2))  
  flatten <- mx.symbol.Flatten(data = pool2)  
  fc1 <- mx.symbol.FullyConnected(data = flatten, num_hidden = 500)  
  tanh3 <- mx.symbol.Activation(data = fc1, act_type = "tanh")  
  fc2 <- mx.symbol.FullyConnected(data = tanh3, num_hidden = num_classes)  
  lenet <- mx.symbol.SoftmaxOutput(data = fc2, name = 'softmax')  
  return(lenet)  
}
```

Fancy structures

- Compute different filters
- Compose one big vector from all of them
- Layer them iteratively




```

def get_symbol(num_classes=1000):
    data = mx.symbol.Variable(name="data")
    # stage 1
    conv1 = ConvFactory(data=data, num_filter=64, kernel=(7, 7), stride=(2, 2), pad=(3, 3), name='1')
    pool1 = mx.symbol.Pooling(data=conv1, kernel=(3, 3), stride=(2, 2), name='pool_1', pool_type='max')
    # stage 2
    conv2red = ConvFactory(data=pool1, num_filter=64, kernel=(1, 1), stride=(1, 1), name='2_red')
    conv2 = ConvFactory(data=conv2red, num_filter=192, kernel=(3, 3), stride=(1, 1), pad=(1, 1), name='2')
    pool2 = mx.symbol.Pooling(data=conv2, kernel=(3, 3), stride=(2, 2), name='pool_2', pool_type='max')
    # stage 3
    in3a = InceptionFactoryA(pool2, 64, 64, 64, 64, 96, "avg", 32, '3a')
    in3b = InceptionFactoryA(in3a, 64, 64, 96, 64, 96, "avg", 64, '3b')
    in3c = InceptionFactoryB(in3b, 128, 160, 64, 96, '3c')
    # stage 4
    in4a = InceptionFactoryA(in3c, 224, 64, 96, 96, 128, "avg", 128, '4a')
    in4b = InceptionFactoryA(in4a, 192, 96, 128, 96, 128, "avg", 128, '4b')
    in4c = InceptionFactoryA(in4b, 160, 128, 160, 128, 160, "avg", 128, '4c')
    in4d = InceptionFactoryA(in4c, 96, 128, 192, 160, 192, "avg", 128, '4d')
    in4e = InceptionFactoryB(in4d, 128, 192, 192, 256, '4e')
    # stage 5
    in5a = InceptionFactoryA(in4e, 352, 192, 320, 160, 224, "avg", 128, '5a')
    in5b = InceptionFactoryA(in5a, 352, 192, 320, 192, 224, "max", 128, '5b')
    # global avg pooling
    avg = mx.symbol.Pooling(data=in5b, kernel=(7, 7), stride=(1, 1), name="global_pool", pool_type='avg')
    # linear classifier
    flatten = mx.symbol.Flatten(data=avg, name='flatten')
    fc1 = mx.symbol.FullyConnected(data=flatten, num_hidden=num_classes, name='fc1')
    softmax = mx.symbol.SoftmaxOutput(data=fc1, name='softmax')
    return softmax

```

Bringing Caffe to MXNet

Caffe is widely used in computer vision

Call Caffe Operators in MXNet

```
import mxnet as mx
data = mx.symbol.Variable('data')
fc1 = mx.symbol.CaffeOp(data_0=data, num_weight=2, prototxt=
    "layer{type:\\"InnerProduct\\" inner_product_param{num_output: 128} }")
act1 = mx.symbol.CaffeOp(data_0=fc1, prototxt="layer{type:\\"TanH\\"}")
fc2 = mx.symbol.CaffeOp(data_0=act1, num_weight=2, prototxt=
    "layer{type:\\"InnerProduct\\" inner_product_param{num_output: 10}}")
mlp = mx.symbol.SoftmaxOutput(data=fc3)
```

Bringing Torch to MXNet



Torch is a popular Lua framework for both scientific computing and deep learning


Tensor Computation

```
import mxnet as mx
x = mx.th.randn(2, 2, ctx=mx.gpu(0))
y = mx.th.abs(x)
print y.asnumpy()
```


Modules (Layers)

```
import mxnet as mx
data = mx.symbol.Variable('data')
fc = mx.symbol.TorchModule(data_0=data, lua_string='nn.Linear(784, 128)', ...
mlp = mx.symbol.TorchModule(data_0=fc, lua_string='nn.LogSoftMax()', ...
```

Distributed Deep Learning



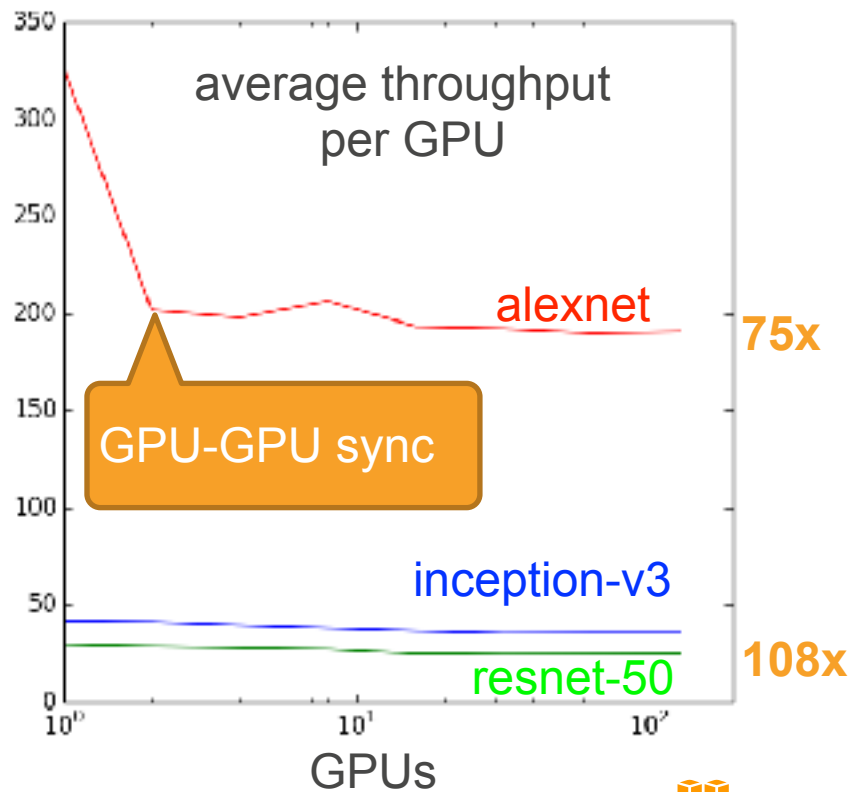
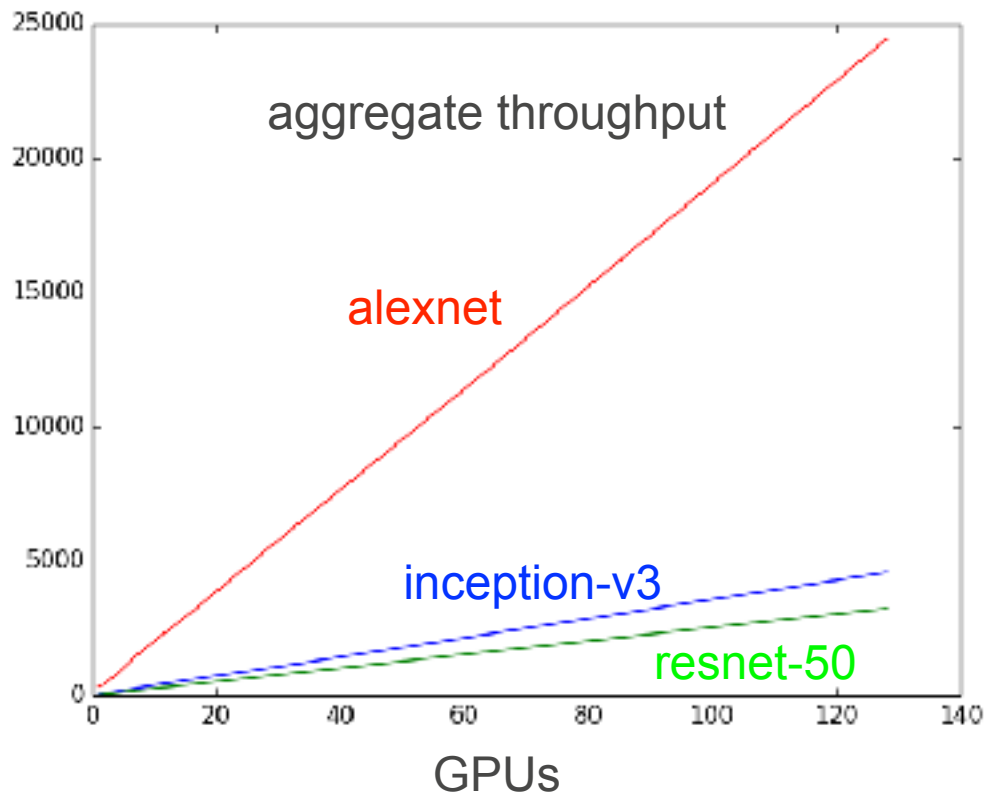
Distributed Deep Learning



2 lines for multi GPU

```
## train
num_gpus = 4
gpus = [mx.gpu(i) for i in range(num_gpus)]
model = mx.model.FeedForward(
    ctx          = gpus,
    symbol       = softmax,
    num_round    = 20,
    learning_rate = 0.01,
    momentum     = 0.9,
    wd           = 0.00001)
model.fit(X = train, eval_data = val, batch_end_callback = mx.callback.Speedometer(batch_size=batch_size))
```

Scaling on p2.16xlarge



AMIs, Cloud Formation and DL Frameworks

- Amazon Machine Images (AMI)
- Deep Learning Frameworks
- Cloud Formation Templates

Amazon Machine Image for Deep Learning

<http://bit.ly/deepami>

- Tool for data scientists and developers
- Setting up a DL system takes (install) time & skill
 - Keep packages up to date and compiled (MXNet, TensorFlow, Caffe, Torch, Theano, Keras)
 - Anaconda, Jupyter, Python 2 and 3
 - **NVIDIA** Drivers for G2 and P2 instances
 - **Intel MKL** Drivers for all other instances (C4, M4, ...)

Getting started

```
acbc32cf4de3:image-classification smola$ ssh ec2-user@54.210.246.140
```

Last login: Fri Nov 11 05:58:58 2016 from 72-21-196-69.amazon.com

Deep Learning AMI for Amazon Linux

This is beta version of the Deep Learning AMI for Amazon Linux.

The README file for the AMI →→→→→→→→→→→→→→→→ /home/ec2-user/src/README.md

Tests for deep learning frameworks →→→→→→→→→→ /home/ec2-user/src/bin

7 package(s) needed for security, out of 75 available

Run "sudo yum update" to apply all updates.

Amazon Linux version 2016.09 is available.

```
[ec2-user@ip-172-31-55-21 ~]$ cd src/
```

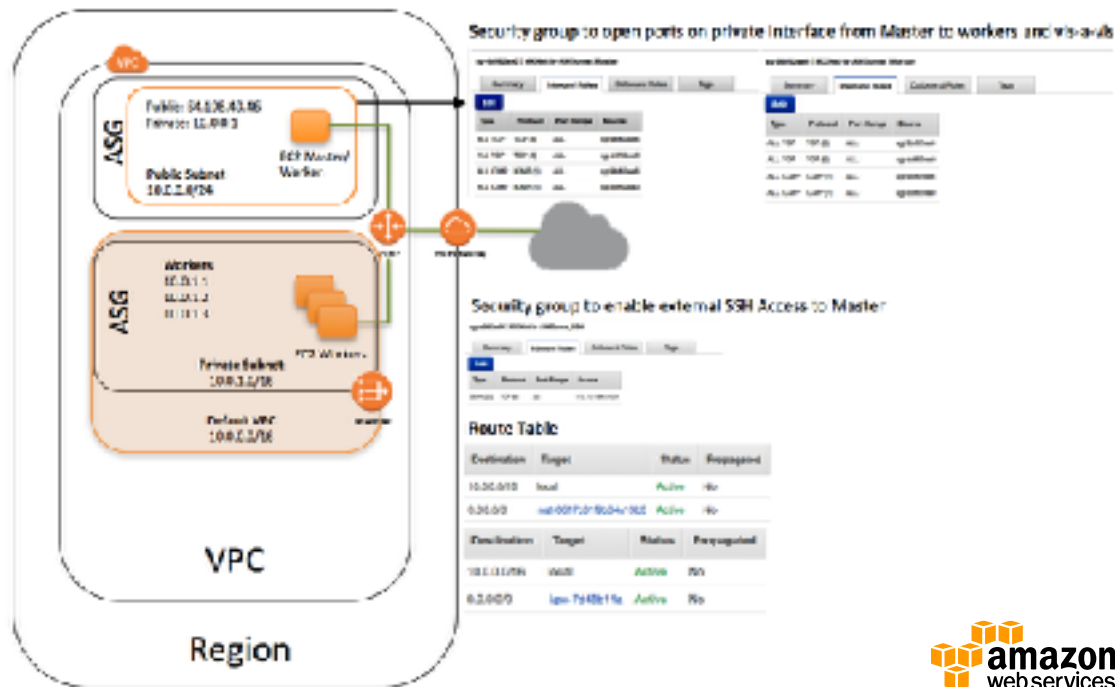
```
[ec2-user@ip-172-31-55-21 src]$ ls
```

anaconda2 bazel caffe cntk keras mxnet OpenBLAS README.md Theano
anaconda3 bin caffe3 demos logs Nvidia_Cloud_EULA.pdf opencv tensorflow torch



AWS CloudFormation Template for Deep Learning

<http://bit.ly/deepcfn>



AWS CloudFormation Components

- **VPC** in the customer account.
- The requested number of **worker instances** in an Auto Scaling group within the VPC. Workers are launched in a **private subnet**.
- **Master instance** in a separate Auto Scaling group that acts as a proxy to enable connectivity to the cluster via **SSH**.
- Two security groups that open ports on the **private subnet** for communication between the master and workers.
- **IAM role** that allows users to access and query Auto Scaling groups and the private IP addresses of the EC2 instances.
- **NAT gateway** used by instances within the VPC to talk to the outside.

Roadmap

- **NNVM Migration** (complete)
- **Apache project** (proposal submitted)
- **Usability**
 - Documentation (installation, native documents, etc.)
 - Tutorials, examples
- **Platform** support
(Linux, Windows, OS X, mobile ...)
- **Language** bindings
(Python, C++, R, Scala, Julia, JavaScript ...)
- **Sparse** datatypes and **LSTM** performance improvements



We are hiring!

{smola, spisakj, mli}@amazon.com

